

Link	Description	Contributor	
<i>Name of the dataset + hyperlink</i>	<i>Brief description of the dataset or how it was acquired</i>	<i>Your Name</i>	
Data1 — SLURP: https://github.com/pswietojanski/slurp	Downloaded human speech recordings and selected utterances matching the project's voice-command labels.	Mark Andrian Macalalad	
Data1 — Fluent Speech Commands: https://zenodo.org/records/11106540	Downloaded spoken commands for lights, music, volume, and temperature, then mapped them to project labels.	Mark Andrian Macalalad	
Data1 — Google Speech Commands v2: https://www.tensorflow.org/datasets/catalog/speech_commands	Used recorded "stop" commands and background-noise samples to supplement command and silence classes.	Mark Andrian Macalalad	
Data2 — eSpeak NG: https://github.com/espeak-ng/espeak-ng	Generated 10,000 recordings: 20 command labels × 10 synthetic voice identities × 10 phrases × five variations in speed, pitch, and volume.	Mark Andrian Macalalad	
Data3 — Chatterbox TTS: https://github.com/resemble-ai/chatterbox ; LibriSpeech: https://www.openslr.org/12/ ; Mozilla Common Voice: https://commonvoice.mozilla.org/	Designed to generate 9,000 synthetic recordings: 20 labels × 30 reference speakers × five phrases × three variants. Each phrase contains 1–2 words. References comprise 24 LibriSpeech and six Filipino-English speakers, with 15 male and 15 female overall. Speaker-disjoint splits: training—20 LibriSpeech + four Filipino-English; validation and test—two LibriSpeech + one Filipino-English each.	Mark Andrian Macalalad	
	The Google Speech Commands Dataset v0.02 is a curated collection of short (approximately one-second) audio recordings of spoken words, specifically designed for training and benchmarking keyword spotting systems. Each recording captures a single spoken command uttered by a diverse set of speakers, making the dataset highly valuable for developing robust, real-world voice-controlled applications. The commands include common terms such as "yes", "no", "up", "down", "left", "right", "on", "off", "stop", and "go", among others. In addition to the primary command recordings, the dataset also provides a set of background noise audio files. These files, stored in a dedicated folder, are intended to support data augmentation techniques and help improve model performance in noisy environments. The dataset has been widely adopted in both academic research and industry applications, serving as a benchmark for lightweight and efficient speech recognition systems.	Joven Daniel Nepomuceno	
Speech commands dataset -- https://www.kaggle.com/code/bhataakib02/echoedgemarvinwithvadfinal	Dataset including one words, we can prune here especially: off, on, stop, etc.	Joven Daniel Nepomuceno	
Synthetic commands speech dataset -- https://www.kaggle.com/datasets/buchner/synthetic-speech-commands-dataset	Synthetic data generated using cosy2. Contains 14k recordings for command 1 & 53k recordings for command 8	Cherry Magdaong	
Synthetic commands speech dataset -- https://github.com/lugan113/SynTTS-Commands-Official	Speaker Verification Corpus FLUENT SPEECH COMMANDS DATASET Dataset contains 97 speakers saying 248 different phrases. The 248 utterances map to 31 unique intents, that are divided into three slots: action, object, and location. The goal in preparing this dataset was to provide a benchmark for end-to-end spoken language understanding models.	Joven Daniel Nepomuceno	
fluent-speech-corpus -- https://www.kaggle.com/datasets/tommyngx/fluent-speech-corpus			
Common Voice by Mozilla -- https://www.kaggle.com/datasets/mozillaorg/common-voice?select=cv-valid-train.csv	General English speech recognition dataset. Has Filipino accent recordings. Around 200k total recordings (multiple English-speaking accents). Timers and Such (Lugosch et al., NeurIPS 2021 Datasets & Benchmarks, arXiv: 2104.01604, license "other-open") has real human recordings for exactly this: SetTimer and SetAlarm intents, mapped to TIMER and ALARM respectively.	King Villasin	
Timers and Such - https://zenodo.org/api/records/4623772/files/timers-and-such-v1.0.zip/content	Only the relevant real subset is downloaded, not the full archive: the source archive is a single 12.2GB zip, and the paper reports only 2,151 real (non-synthetic) utterances across all 4 of its intents combined. So I created a script to only take the ones that match the schema we're using.	Quiel Quiwa	
Snips SLU - https://huggingface.co/datasets/MWilinski/snips_sl_u_v1.0	It contains recordings of commands for smart home appliances in English, with info about c	Quiel Quiwa	